WHITEPAPER VERSION 1.0

The Vibegraph™

Your Vibes, Codified.

An open framework for codifying personal and brand identity as portable context for AI tools and agents.

July 11, 2026 vibegraph.ai

Ryan Charleston vibegraph.md

Introduction

Every AI tool you use starts from zero. It doesn't know your voice, your values, your goals, or your taste. So it guesses, and it guesses toward the statistical average. The result is output that could belong to anyone, which is another way of saying it belongs to no one.

The vibegraph fixes the input instead of endlessly editing the output. It is a structured, portable, human-owned system that codifies who you are: your personality, your purpose, your brand positioning, and your visual identity, plus optional modules covering the working context of your life and/or business. Any AI tool, agent, or memory system can consume it. None of them own it.

Two design decisions separate the vibegraph from the wave of AI memory products. First, it is grounded in established frameworks rather than passive data collection: the Big Five "OCEAN" personality model, Ikigai for purpose and meaning, and industry best practices and standards in brand strategy for positioning and aesthetics (with organizational equivalents for business vibegraphs, drawn from Aaker's brand personality dimensions, Jungian brand archetypes, and the Golden Circle). Second, it works today, with zero platform adoption required. Paste a vibegraph into Claude, ChatGPT, Grok, or Gemini and the next output sounds like you.

> *Memory layers remember what happened. A vibegraph defines who you are.*

This paper describes the problem, the architecture, the personal and business variants of the framework, how AI systems consume a vibegraph, the security model such a concept demands, and the open specification at vibegraph.md.

---

1. The Problem: AI Doesn't Know Who You Are

Large language models are trained on the written output of millions of people and optimized to produce the most probable response. Probable means average. When a model knows nothing about you, the average is the best it can do.

Anyone who uses AI for real work has felt the consequences:

You repeat yourself constantly. Every new chat, every new tool, every new agent begins with the same ritual: here's who I am, here's what I do, here's how I write, here's what I'm working on. The context you type into one tool is gone the moment you open another.

The output is generic. "AI slop" became a household term for a reason. Ask ten professionals in the same field to generate a LinkedIn post on the same topic and the results are nearly interchangeable, because the model is drawing on the same averaged prior for all ten. The people who complain loudest are exactly the people with the most to lose: creators, coaches, consultants, and founders whose personal brand is their business. Their differentiation is their voice or unique “vibes”, and generic output erodes that brand with every post.

Platform memory doesn't solve this, and creates a new problem. ChatGPT remembers some things about you. Claude remembers other things. Neither can see the other's memory, you can't fully inspect either one, and none of it transfers when a better tool ships next quarter. Memory features are useful, but they are fragmented, passive, and platform-locked. What they accumulate is a history of what you did, filtered through what a model happened to notice. That is not the same as a deliberate statement of who you are, today, nor does it keep track of who you are becoming.

Agents raise the stakes. A chatbot that misreads your tone wastes a prompt. An autonomous agent that misreads your intent takes actions on your behalf: it sends the email, books the meeting, publishes the post. The gap between “AI that knows you” and “AI that guesses” stops being a quality issue and becomes a trust issue.

The software industry has already solved a version of this problem for code. AGENTS.md, an open markdown convention that tells coding agents how a project works, spread to more than 60,000 repositories in about a year because the payoff was immediate: write the file once, and every agent that touches your codebase gets better on the next task. No committee, no SDK, no permission from anyone.

There is no equivalent for a human being. That is the gap the vibegraph fills.

---

2. What Is a Vibegraph?

A vibegraph is a structured collection of documents that codifies a person's (or organization's) identity and working context in a format AI systems can consume. Plain markdown at its simplest. A permissioned, encrypted data vault at its most complete. You own it, you host it where you choose, and you decide which parts any given tool or agent can see.

The name is literal. Your “vibes” — personality, taste, voice, values, purpose, aesthetic — are the things AI gets wrong about you by default, because they were never written down anywhere a machine could read them. The vibegraph writes them down and keeps them up-to-date.

Design Principles

Human-authored, not scraped. A vibegraph is built deliberately, through structured self-assessment and brand-strategy work, not inferred from your inbox. Passive inference produces a model's opinion of you. A vibegraph is your statement of record. (The two can coexist; a vibegraph makes an excellent seed and correction layer for passive memory systems.)

Useful on day one, with zero adoption required. This is the load-bearing principle. A vibegraph delivers its value the moment its owner pastes it into any AI tool, LLM, or agent. Although welcomed, it's not dependent on third-party platforms to support it as a standard. Proposed web standards that depended on platform adoption have a poor track record; the ones that win give individual users an immediate payoff with tools they already have and use. The vibegraph is designed for the second pattern.

Portable and sovereign. The vibegraph is a framework and system that can be exported to a document/file format. It moves with you across tools, models, and vendors. If a platform shuts down or a better one launches, your vibegraph comes along unchanged.

Grounded in established frameworks. The identity layer is not improvised. It uses instruments and frameworks with decades of practice behind them, so the output is structured, comparable, and complete rather than a freeform “about me” essay.

Private by default, shared by choice. Every element in a vibegraph carries its own visibility setting. The default is private and local. Section 7 covers this in depth.

Extensible. A minimal vibegraph is just the identity core: a few documents. From there, owners add modules for whatever domains matter to them, at whatever depth they choose and feel comfortable documenting.

Glossary

Vibegraph — Your vibes, codified.
A portable, structured identity and context framework that codifies personality, purpose, and brand so AI tools and agents act in alignment with their owner's goals, voice, and taste. A vibegraph can represent a person or an organization.

Core Identity — Who you are.
The foundational half of a vibegraph: personality, purpose, and brand (positioning plus aesthetics). Small, stable, and safe to share with any AI tool.

Knowledge Base — What you know and how you operate.
The extensible half of a vibegraph: self-defined modules containing the documents, data, and working context of your life or business. Creators often call this their second brain; businesses call it a knowledge base or wiki. Modules are gated and permissioned based on user preferences.

Vibeclone — You, digitally remastered.
An AI agent powered by a vibegraph that thinks, writes, and acts in its owner's likeness. Covered briefly in Section 9; a dedicated paper will follow.

3. Architecture: Two Parts

Every vibegraph, personal or business, has the same two-part structure.

The **Core Identity** answers the question every AI tool silently asks and never gets answered: who is this person? It contains the personality profile, the purpose work, and the brand layer. It is deliberately small — the whole core fits comfortably inside a single model context window — because it is meant to travel everywhere. It is the part you paste into a chat, attach to a project, or expose to an agent without a second thought, because it contains nothing you wouldn't put on a well-written about page (unless you choose to go deeper).

The **Knowledge Base** answers the follow-up question: what is this person working with? It is a set of self-defined modules — health, finances, relationships, career, notes, time management, skills, goals, or anything else the owner wants to include — each containing documents, files, data, and protocols. The Knowledge Base is where depth lives, and it is also where sensitivity lives, which is why every module and every item within it carries its own access controls.

The split matters for three reasons.

First, it matches how AI consumption actually works. Identity context belongs in every interaction; domain context belongs only in relevant ones. A model writing your newsletter needs your voice and positioning, not your medical history. Separating the always-relevant from the sometimes-relevant keeps context windows lean and permission decisions simple.

Second, it matches how trust works. The Core is shareable by design. The Knowledge Base is gated by design. A clean architectural line between the two is easier to secure and easier to reason about than a single blob with per-paragraph exceptions.

Third, it scales in both directions. A minimal vibegraph — Core only — is genuinely useful and can be built quickly with proper tools. A maximal vibegraph, with a dozen modules maintained over years, approaches something like a personal operating system. Both are valid. The framework doesn't force anyone up the curve.

4. The Personal Vibegraph

4.1 Core Identity

The personal Core Identity has three layers, built in order, because each one feeds the next.

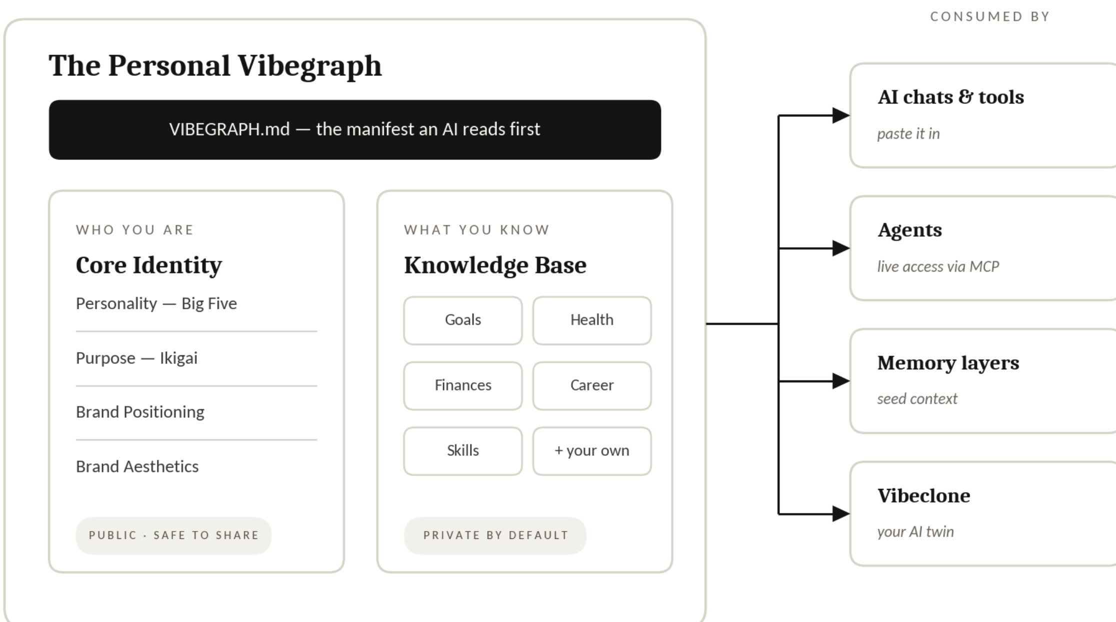

Figure 1 — The personal vibegraph: a manifest, a Core Identity built on Big Five and Ikigai, and a permissioned Knowledge Base of life-domain modules — consumed everywhere from a chat window to a memory layer.

Personality: the Big Five (OCEAN). The foundation is a Big Five personality assessment measuring Openness, Conscientiousness, Extraversion, Agreeableness, and Neuroticism. The Big Five is used rather than more popular typology instruments for a simple reason: it is the model with the strongest empirical support in personality psychology, it produces dimensional scores rather than type labels, and dimensional data is exactly what a machine consumer can use. “High openness, moderate extraversion, low neuroticism” gives a model calibration that “ENFP” does not. The vibegraph stores the scores, the facet breakdown, and a written interpretation in the owner's own words.

Purpose & Meaning: Ikigai. The second layer codifies meaning and direction using Ikigai 2.0, the corrected and expanded treatment of the Japanese concept of ikigai (developed by Kyle Kowalski at Sloww). Where the pop-culture ikigai Venn diagram reduces purpose to a career sweet spot, Ikigai 2.0 treats purpose, meaning, and vocation as related but distinct, and produces a richer artifact: what gets you up in the morning, what you find meaningful, what you're building toward, and why. For an AI consumer, this layer is what turns “write a post about productivity” into a post about productivity *that connects to what its author actually cares about.*

Brand: positioning and aesthetics. The third layer is the personal brand, in two halves. *Positioning* covers the strategic side: positioning principles, brand name, unique value proposition, purpose-vision-mission, core values, tone and voice, messaging and narratives, brand keywords, short and long bios, achievements, interests, inspirations, and online presence. *Aesthetics* covers the visual side, and it produces concrete artifacts rather than vague direction: hex color values, typography pairings, logo and symbol options, iconography, imagery guidelines, and photography direction. When an AI tool generates a slide, a thumbnail, or a landing page for you, this is the layer that keeps it on-brand without a design review.

Together, the three layers give any AI system what a great ghostwriter, a great designer, and a great strategist would each need a dozen sessions to absorb.

4.2 Knowledge Base

The Knowledge Base is organized into modules. The framework suggests a starting taxonomy; owners rename, merge, split, and invent modules freely. The default suggested personal set:

Module	Typical contents
Health & Wellness	Sleep, nutrition, fitness, habits and rituals, health records and documents
Relationships	Family, friends, coworkers, partners and clients, audience
Finances	Budget, savings and investments, wealth and tax documents
Career	Résumé/CV and career history, current role, current business info
Notes & Ideas	Quick notes, brain dumps, bookmarks and read-later queues, journals
Time Management	Principles, routines and schedules, planner, calendar conventions
Skills	Methods, protocols, skills written for AI consumption, reference docs
Tech Stack	Tools and hardware/software, current/intended use, access and privileges
Goals	Short- and long-term, with daily/weekly/monthly/yearly cadence

A note on method: readers who maintain a second brain will recognize that this is a life-domain taxonomy, closer to Johnny Decimal or the “Areas” of PARA than to Zettelkasten's emergent linking. That is deliberate. Systems optimized for human creative recall reward loose association; systems optimized for machine retrieval reward stable, addressable domains an agent can be granted access to, one at a time. The vibegraph does not replace your PKM practice. It sits in front of it as the permissioned, machine-facing layer, and it can happily link out to an Obsidian vault, Notion workspace, or a folder of markdown files.

The most common failure mode of personal knowledge systems is well documented: capture everything, retrieve nothing, abandon the system. The vibegraph's defense is the Core/module split itself. The Core alone delivers the headline value. Modules are added when a real use case demands one, not because an empty template exists.

5. The Business Vibegraph

Organizations need the same thing individuals do: a portable, structured identity that keeps every AI touchpoint on-brand and in-context. But an organization is not a person, and the frameworks do not transfer one-to-one. A company cannot take a personality test, and ikigai does not describe a business. The business vibegraph keeps the two-part architecture and swaps the instruments.

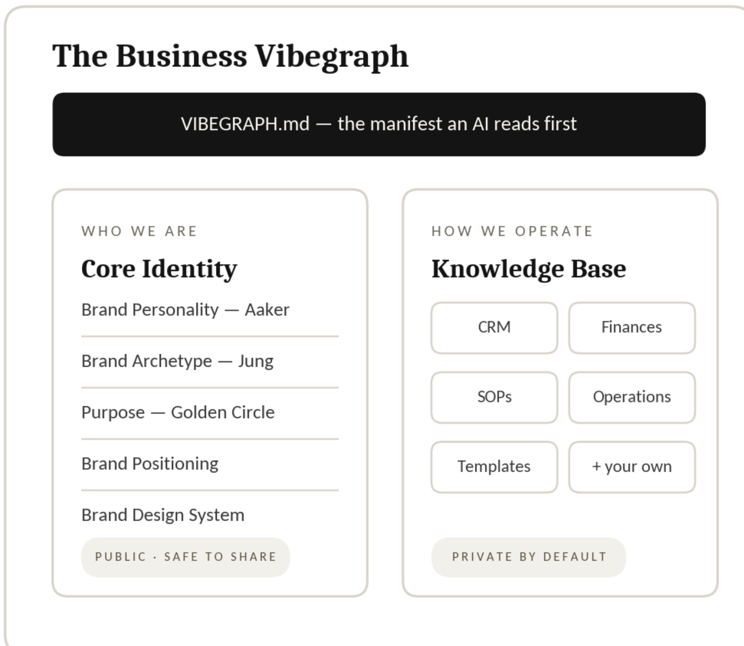

Figure 2 — The business vibegraph: the same two-part architecture with the instruments swapped for organizational equivalents — and the modules restructured around repeatability.

5.1 Business Core Identity

Brand personality: Aaker's five dimensions. In place of the Big Five, the business Core uses Jennifer Aaker's brand personality framework — Sincerity, Excitement, Competence, Sophistication, Ruggedness — the closest thing brand research has to an empirically derived trait model for organizations. Like OCEAN, it produces dimensional, machine-usable calibration rather than a slogan.

Brand character: the twelve archetypes. Aaker's dimensions describe a brand's personality; a Jungian archetype (Hero, Sage, Creator, Outlaw, Caregiver, and the rest) gives it a coherent character to write and design from. In practice the two work together: choose the archetype, then use the dimensions to describe and measure how it shows up. Both live in the business Core.

Purpose: the Golden Circle, mission, vision, values. In place of Ikigai 2.0, the business Core codifies purpose through Simon Sinek's Why/How/What structure alongside conventional mission, vision, and values statements. (Small businesses running on EOS can drop their Vision/Traction Organizer in here nearly unchanged.) Usefully, Sinek designed the Golden Circle to apply to individuals as well as organizations, which gives personal and business vibegraphs a shared spine: a founder's personal Why and their company's Why can be written in the same shape and checked against each other.

Brand: positioning and design. The brand layer carries over from the personal framework with one adjustment in emphasis: the aesthetics layer becomes a full design system (design principles, symbols and logos, palette, typography, iconography, imagery, and templates for the assets a business produces repeatedly — social posts, decks, email signatures, landing pages, media kits).

5.2 Business Knowledge Base

The business modules shift from life domains to operating systems — the things that make delegation and consistency possible:

Module	Typical contents
Relationships (CRM)	Employees, advisors, investors, partners, vendors, customers, audience
Finances	Budget, P&L, payroll, investments, valuation
Operations	Business entity, HR, legal
Skills	Playbooks, SOPs, workflows, skills written for AI, reference docs
Management	Principles, routines, planning, calendar systems
Brand Templates	Post templates, decks, email templates, lead magnets, landing pages, media kits
Goals	Short- and long-term, with cadence
Notes & Ideas	Capture, bookmarks, journals/logs
Tech Stack	Tools, access and privilege protocols

The structural difference from the personal Knowledge Base is worth stating plainly, because it explains why the module lists diverge. A personal knowledge system optimizes for one person's recall and creative output. A business knowledge system optimizes for *repeatability*: any operator — teammate, contractor, or agent — following the same playbook should produce the same result. That is why SOPs and templates, which barely appear in personal systems, are the center of gravity in business ones. It is also why the business vibegraph is unusually well suited to agents: an SOP written clearly enough for a new hire is most of the way to being a protocol an agent can execute.

5.3 What maps, what doesn't

Personal	Business	Relationship
Big Five (OCEAN)	Aaker's five dimensions + archetype	Replaced: same job (dimensional personality), different instrument
Ikigai 2.0	Golden Circle + mission/vision/values	Replaced: same job (purpose), shared Why structure
Brand positioning	Brand positioning	Carries over nearly intact
Brand aesthetics	Brand design system + templates	Carries over, deepens
Life-domain modules	Operating-system modules	Restructured around repeatability

A solopreneur, usefully, needs both — and the overlap is small enough that maintaining a personal vibegraph and a lightweight business vibegraph side by side is practical. The personal one governs voice and identity; the business one governs the machine that sells and delivers.

---

6. How AI Systems Consume a Vibegraph

A vibegraph is only as good as the ways it can be used. There are four, in increasing order of sophistication, and the first one requires nothing from anyone.

1. Direct context. Paste the Core Identity into any chat, or attach it to any tool that accepts files. This works today, in every AI product on the market, and it is the canonical consumption mode. For AI builders, the corollary: a vibegraph is structured markdown, so if your product accepts text, it already supports vibegraphs. The feedback loop is immediate: the very next response is calibrated to its owner. This is the same loop that drove AGENTS.md through the developer ecosystem — write the file once, watch every tool get better. VIBEGRAPH.md applies the same idea to a person instead of a codebase.

2. Persistent workspace context. Most serious AI tools now have a persistent-context surface: Claude's Projects, ChatGPT's custom GPTs and project instructions, and their equivalents elsewhere. A vibegraph exports cleanly into all of them. Set it once per workspace and every conversation in that workspace starts calibrated.

3. Live access over MCP. The Model Context Protocol gives agents a standard way to request data from external sources at runtime. A vibegraph served over MCP lets an agent ask for exactly the module it needs, when it needs it, subject to the owner's permissions — your writing agent reads Core Identity and the Brand Templates module; your finance agent reads Finances and nothing else. This is where the per-item permission model (Section 7) does its real work, and it is the consumption mode that turns a vibegraph from a document into critical infrastructure.

4. Seeding memory layers. AI memory systems — the platform-native ones and the dedicated memory infrastructure used by agent builders — all face a cold-start problem: they know nothing until they've watched you for weeks, and what they learn is inference, not fact. A vibegraph solves this cold start problem. Loaded as seed context, it gives a memory layer a verified, owner-authored foundation that passive observation can then extend. This is the intended relationship between the vibegraph and the well-funded memory category: they are the substrate, the vibegraph is the schema. Memory layers remember what happened. A vibegraph defines who you are. The two compose.

One relationship worth making explicit: the vibegraph is to a person what AGENTS.md is to a repository. Same medium (plain, human-readable markdown), same philosophy (a predictable place for the context machines need), same adoption physics (individual users get instant value; ecosystem support compounds it but is never required).

> *Where AGENTS.md tells an agent how to work on your code, a vibegraph tells an agent how to work as, and for, you.*

---

7. Security and Privacy

A complete vibegraph is a concentrated dossier: personality profile, purpose, finances, health, relationships, goals. In the wrong hands it is a toolkit for impersonation, social engineering, and targeted fraud. Any honest presentation of this framework has to treat security as architecture, not as a policy page. This section describes the threat model and the design that answers it.

7.1 Threat Model

The realistic attack surfaces, roughly in order of likelihood:

Compromised or over-permissioned AI tool access. The newest and most distinctive risk. Prompt-injection attacks that exfiltrate data through AI tool integrations moved from theory to documented incidents in 2025, including zero-click exfiltration through enterprise AI assistants and data theft through poisoned MCP tool descriptions. The security community's consensus position is sobering: prompt injection cannot currently be fully eliminated, only contained. Any system that exposes personal data to AI tools must therefore assume that some tool, someday, will be manipulated, and must limit what that tool can access in the first place.

Server breach. If vibegraphs are stored on a provider's servers in readable form, the provider is a honeypot. One breach exposes every customer's most sensitive self-description at once.

Account takeover. Credential stuffing and phishing against individual accounts.

Insider access. Anyone at a hosting provider who can read customer data is a risk, however trustworthy the company.

Training-data leakage. Personal context sent to model providers without contractual protection may end up in training corpora.

7.2 Design Response

Private and local by default. The reference posture for a vibegraph is files on hardware the owner controls. The open framework at vibegraph.md assumes local-first operation: your vibegraph lives on your machine, and only the elements you deliberately expose ever leave it. Owners who self-host accept responsibility for their own device security, and the framework documentation says so plainly rather than pretending a local folder is a vault.

Per-item permissions, deny by default. Every module, and every item within a module, carries a visibility setting: private (never exposed), scoped (exposed to named tools or agents, for named purposes), or public (freely shareable, like the Core). Nothing is exposed by omission. The Core Identity is designed to be safe at "public" for most owners; the Knowledge Base defaults everything to "private."

Scoped, logged agent access. When a vibegraph is served to agents over MCP, access is granted per-agent and per-module, never as a blanket. The full graph is never exposed as a single surface. All agent reads are logged, so the owner can always answer the question "what has this tool actually seen?" Given the prompt-injection reality described above, tight scoping is not paranoia; it is the mitigation. An agent that was never granted the Finances module cannot leak it, no matter how thoroughly the agent itself is compromised.

Zero-knowledge encryption for hosted vibegraphs. For owners who prefer a managed service (the hosted app at vibegraph.ai), the storage design follows the model proven by password managers and encrypted-mail providers: client-side encryption, with keys derived on the owner's device and never transmitted. The provider stores ciphertext it cannot read. This carries real trade-offs that deserve an honest statement: a zero-knowledge provider cannot search your data server-side, cannot recover it if you lose your keys, and cannot hand over readable data under compulsion, because it never had any. For a document this sensitive, those trade-offs are the point.

No-train inference, with escalation paths. AI features that process vibegraph content should run against APIs whose terms exclude customer data from model training — the standard posture of major providers' commercial APIs, and the pragmatic baseline for most owners. Owners with stricter requirements can point the framework at local models or privacy-focused inference providers like venice.ai; the vibegraph, being a file format rather than a service, does not care which model reads it.

Future consideration of verifiable credentials. The W3C's Verifiable Credentials standard (v2.0, a formal Recommendation since May 2025) and selective-disclosure techniques point toward a future in which a vibegraph can *prove* an attribute — a credential, an age, a verified identity — without revealing the underlying data. That is a genuinely useful future for a portable identity layer, and the framework is designed not to preclude it. It is deliberately excluded from this version of the vibegraph whitepaper, along with blockchain storage of any kind, because shipping proven encryption today beats promising novel cryptography tomorrow.

7.3 The Risk Summary

A vibegraph concentrates risk in exchange for concentrating value; the design's job is to keep the exchange favorable. Local by default, deny by default, scope everything, encrypt what's hosted so the host can't read it, and log every agent access. No security section can promise safety. This one promises that the sensitivity of the data was the first architectural constraint, not the last.

---

8. Use Cases

High-fidelity AI output. The founding use case. A creator, marketer, or founder with a vibegraph gets output in their voice, aligned with their positioning, styled to their brand, from any AI tool, on the first generation instead of the fifth revision. The people this serves most are those whose personal brand is the business: solopreneurs, founders, coaches, consultants, content creators, educators, agency founders, fractional executives.

Consistency across the whole stack. The same Core Identity feeds the chat tool, the writing assistant, the design generator, the email drafter, and the agent that queues social posts. One source of truth ends the drift between tools that each hold a slightly different, slightly stale picture of you.

Agency client onboarding, upgraded. Marketing and branding agencies extract a shallow vibegraph from every client today; they just call it a discovery questionnaire, it lives in a Google Doc, and it gets read twice. Rebuilt as a vibegraph, client discovery produces a living, structured client-context file the agency loads into every AI tool it uses for that account, for the life of the engagement. Onboarding output stops being paperwork and becomes reusable infrastructure. For agencies, one intake process yields a per-client asset; for clients, the vibegraph is theirs to keep when the engagement ends.

Seeding your digital twin. A vibegraph is the natural training substrate for an AI clone of oneself (see Section 9) and the natural seed for any personal-AI or memory product, giving passive systems a verified foundation instead of weeks of cold-start inference.

The résumé, replaced. A résumé is a one-dimensional, self-flattering PDF. A shared slice of a vibegraph — Core Identity plus a Career module — gives an employer or client a structured, honest, machine-readable picture of how a person thinks, works, and communicates. Whether hiring processes are ready for this is an open question; the artifact, at least, is finally made possible by the vibegraph.

Forms, applications, and intake, automated. Medical intake, loan applications, vendor onboarding, conference speaker forms: tedious, repetitive, and answerable almost entirely from a well-maintained vibegraph. An agent with scoped access to the relevant modules fills them in seconds and asks the owner only for what's genuinely new. Vibegraph owners save a significant amount of time on the task of filling out forms.

9. The Vibeclone

In the future, the most ambitious consumer of a vibegraph will be a Vibeclone: an AI agent that runs on a person's vibegraph and operates in their likeness — thinking through their frameworks, writing in their voice, making the calls they would make, within the permissions they've granted.

The distinction between a vibeclone and today's “AI clone” products is the direction of service. Existing clone platforms point outward at an audience: fans and clients talk to a chatbot trained on a creator's content. A vibeclone points inward, working *for* its owner across their actual tools and tasks, with the vibegraph as its persistent, owner-authored source of self. The better the vibegraph, the less the clone has to guess.

The vibeclone raises its own questions — capability boundaries, disclosure norms, delegation limits, identity verification — that deserve more than a subsection. A dedicated Vibeclone paper (forthcoming) will address and answer these questions. For the purposes of this whitepaper, one sentence suffices: the vibegraph is the prerequisite. There is no faithful clone of an uncodified person.

10. The Open Framework & Application

The vibegraph ships in two forms, deliberately.

The Open Framework: The open specification can be found at vibegraph.md. The framework itself — the document structure, the module taxonomy, the schema, and starter templates — is published free under a permissive license, in the same spirit (and the same medium) as AGENTS.md. Anyone can build a vibegraph by hand, with their own tools, their own storage, and their own security posture. Developers and tinkerers can extend it, wrap it, serve it over MCP, and build products that consume it, without asking permission. An identity standard that isn't open isn't a standard; it's a lock-in scheme with a manifesto.

The Application: The guided application can be found at vibegraph.ai. Building a vibegraph by hand is real work. It takes time and commitment. A proper personality assessment, honest purpose work, and brand strategy development historically could cost thousands (if not tens of thousands) of dollars with a human brand strategist. The hosted vibegraph application packages this work into a convenient tool at a fraction of that cost. The app offers a guided, AI-assisted vibegraph building process — free to start with the personality assessment, a single one-time unlock for the full Core Identity build — plus optional ongoing maintenance, monitoring, agent connections, and the zero-knowledge hosting (described in Section 7) for those who want it. The vibegraph app is a convenience layered on the open framework, never a gate in front of it.

11. Summary

The last two years (mid 2024 – mid 2026) settled the question of whether AI can produce competent work. The next two will be about context and whose work it produces. Left uncalibrated, every model regresses to the same mean, average output or "slop", and everyone who relies on it sounds a little more like everyone else. Your unique personal moat in the agentic era will rely on a well-defined vibegraph and the personal brand that exists within it.

The fix is not a smarter model. It is better input and context: a deliberate, structured, portable statement of identity that any AI system can read and no AI platform can own. Small enough to paste into a chat today. Extensible enough to run an agent fleet tomorrow. Private by default, permissioned by design, and grounded in frameworks older and sturdier than any of the tools that will consume it.

> *The Vibegraph is your vibes, codified.*

Visit **vibegraph.md** to download the open framework or **vibegraph.ai** to start building your vibegraph with the app.

References

The vibegraph stands on established frameworks and open standards. The works and authors below inform or are embedded in its architecture.

Identity & Psychology

1. McCrae, R. R., & Costa, P. T. (1987). "Validation of the Five-Factor Model of Personality Across Instruments and Observers." *Journal of Personality and Social Psychology*, 52(1), 81–90. The foundational validation of the Big Five (OCEAN) model.
2. Goldberg, L. R. (1993). "The Structure of Phenotypic Personality Traits." *American Psychologist*, 48(1), 26–34.
3. International Personality Item Pool (IPIP) — public-domain Big Five assessment instruments. ipip.ori.org
4. Kowalski, K. *Ikigai 2.0: Evolving the Ikigai Diagram for Life Purpose* — the corrected and expanded ikigai framework used in the personal Core Identity. Sloww. sloww.co/ikigai-2-0

Brand Strategy

5. Aaker, J. L. (1997). "Dimensions of Brand Personality." *Journal of Marketing Research*, 34(3), 347–356. The five-dimension brand personality framework used in the business Core Identity.
6. Jung, C. G. (1959). *The Archetypes and the Collective Unconscious.* Princeton University Press. The psychological foundation of archetype theory.

7. Mark, M., & Pearson, C. S. (2001). *The Hero and the Outlaw: Building Extraordinary Brands Through the Power of Archetypes.* McGraw-Hill. The application of Jung's twelve archetypes to brand building.

8. Sinek, S. (2009). *Start with Why: How Great Leaders Inspire Everyone to Take Action.* Portfolio. The Golden Circle (Why/How/What) framework. simonsinek.com

9. Wickman, G. (2011). *Traction: Get a Grip on Your Business.* BenBella Books. The Entrepreneurial Operating System (EOS) and Vision/Traction Organizer. eosworldwide.com

Knowledge Management

10. Forte, T. (2022). *Building a Second Brain.* Atria Books. The PARA method and the modern "second brain" practice. buildingasecondbrain.com

11. Ahrens, S. (2017). *How to Take Smart Notes.* The Zettelkasten method, after the practice of sociologist Niklas Luhmann.

12. Noble, J. *The Johnny.Decimal System* — a numbered life-domain taxonomy for organizing information. johnnydecimal.com

AI Standards & Protocols

13. AGENTS.md — an open markdown convention for guiding AI coding agents, stewarded by the Agentic AI Foundation under the Linux Foundation. agents.md

14. Model Context Protocol (MCP) — an open protocol for connecting AI applications to external data sources and tools, originated by Anthropic and hosted by the Linux Foundation. modelcontextprotocol.io

Security & Privacy

15. W3C (2025). *Verifiable Credentials Data Model v2.0.* W3C Recommendation, May 15, 2025. w3.org/TR/vc-data-model-2.0

16. OWASP Foundation. *OWASP Top 10 for Large Language Model Applications* — the reference taxonomy for prompt-injection and related LLM security risks. genai.owasp.org

17. 1Password. *1Password Security Design* white paper — the zero-knowledge, client-side encryption model referenced in Section 7. 1password.com/security